

② Obtain dataset of stress-ng performance vectors for each machine, as well as the performance vector corresponding to the application.




machine	af-alg	atomic	bigheap	brk	bsearch	cache
c220g1.quiho.schedock-PG0.wisc.ccloudlab.us	7460.525535	680513.700483	37349.785746	692268.660998	369.432133	3.566657
c220g2.quiho.schedock-PG0.wisc.ccloudlab.us	7403.865310	682288.252459	36636.804958	682496.628441	370.318982	3.699981
c6220.quiho.schedock-PG0.apr.emulab.net	11054.153981	564287.762689	11611.122370	303527.165342	306.383069	4.299996
c6320.quiho.schedock-PG0.clemson.ccloudlab.us	7757.281682	451002.732744	26667.801230	506820.492253	251.355372	3.066660
c8220.quiho.schedock-PG0.clemson.ccloudlab.us	7686.662436	450701.333302	27311.612889	511324.124936	251.630304	2.933328
d2100.quiho.Schedock.emulab.net	4379.832362	164771.809946	2025.253896	32487.221983	173.951188	41.466890
d430.quiho.Schedock.emulab.net	4364.033699	164629.906329	1907.548357	34439.766613	176.327592	38.000142
d530.quiho.Schedock.emulab.net	4350.648073	164693.273432	2005.794649	32321.318362	175.592944	39.000188
d710.quiho.Schedock.emulab.net	4356.095055	164447.914798	1915.892418	33937.885275	173.925279	40.566921
d820.quiho.Schedock.emulab.net	4357.132861	164863.654692	1936.961068	31408.002412	173.915386	37.600261
d360.quiho.emulab-net.utahddc.geniracks.net	6661.338910	433914.302217	26323.321728	478416.274642	250.914831	3.333331
dwill.soe.ucsc.edu	13223.031682	657467.782752	43978.646205	691672.362863	371.884756	11.266639
lssdm-41	22839.765193	523305.564038	6025.801038	208221.311786	173.388412	19.533220
m510.quiho.schedock-PG0.utah.ccloudlab.us	7507.446956	531755.228881	46000.484393	767875.363427	299.498324	5.766648
pc2400.quiho.emulab-net.uky.emulab.net	9613.824407	380447.538905	9341.089105	175586.465802	236.164213	11.966682
pc3000.quiho.Schedock.emulab.net	4361.731547	164690.357913	2422.164826	37981.053613	176.258683	37.933498
pc3300.quiho.emulab-net.uky.emulab.net	7670.092317	409655.727577	9621.317555	177406.858332	229.030712	14.800006
pc3400.quiho.emulab-net.uky.emulab.net	7802.481602	409704.016410	11132.209174	180905.497098	229.222454	14.800010
pc3500.quiho.emulab-net.uky.emulab.net	7650.021945	409757.666632	10076.285587	189780.573050	228.243323	14.238500
r320.quiho.schedock-PG0.apr.emulab.net	11054.577959	559096.754508	10002.069980	226115.721011	311.061257	3.966669
r720.quiho.schedock-PG0.apr.emulab.net	11148.859688	560523.049680	13725.433750	334987.529508	308.742572	4.100001
scruffy.soe.ucsc.edu	7841.064578	605547.922998	19078.094634	382625.637476	250.386584	6.099999

X


y



③ Create a prediction model via sklearn's gradient boosting. This produces a tree ensemble.



① Execute the battery of microbenchmarks (stress-ng), as well as the application being profiled, on all the available machines.



④ Create the IRUP by obtaining the relative ranking of feature importance out of the ensemble.

